


Flexible Instruction Processors

Wayne Luk & Peter Y.K. Cheung
Imperial College
CASES 2000



“Not So General” Purpose Processors

- ◆ GPPs lose performance when dealing with non-standard operations and data
- ◆ Particularly apparent in embedded systems
 - Cell phones
 - Medical appliances
 - Digital cameras
 - PDAs

Two roads to customization

◆ Augment GPPs with programmable logic

- Couple standard processor (ARM, MIPS) with an FPGA fabric
- Fixed processor instruction set
- FPGA implements custom instructions

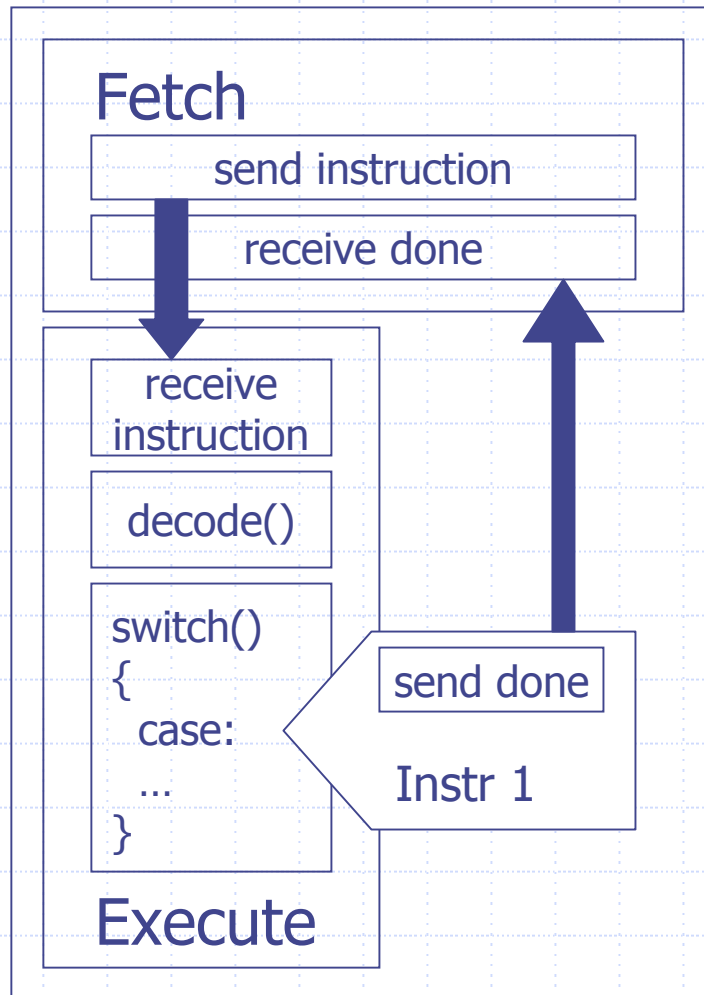
◆ Implement them in FPGAs

- Customize instructions at compile time or at run time

Flexible Instruction Processors


- ◆ FIPs are assembled by parameterized processor templates
- ◆ Allows customized hardware to be exposed as new instructions
- ◆ Critical resources can be adjusted according to domain requirements
- ◆ Can implement same instruction set with different design tradeoffs

FIPs Template



- ◆ Templates are modules connected together by communicating channels
- ◆ Can instantiate different styles of processors
 - Stack, register, etc.
- ◆ Customizations
 - Removal of resources
 - Changing data widths
 - Vary degree of pipelining
- ◆ Implemented in Handel-C

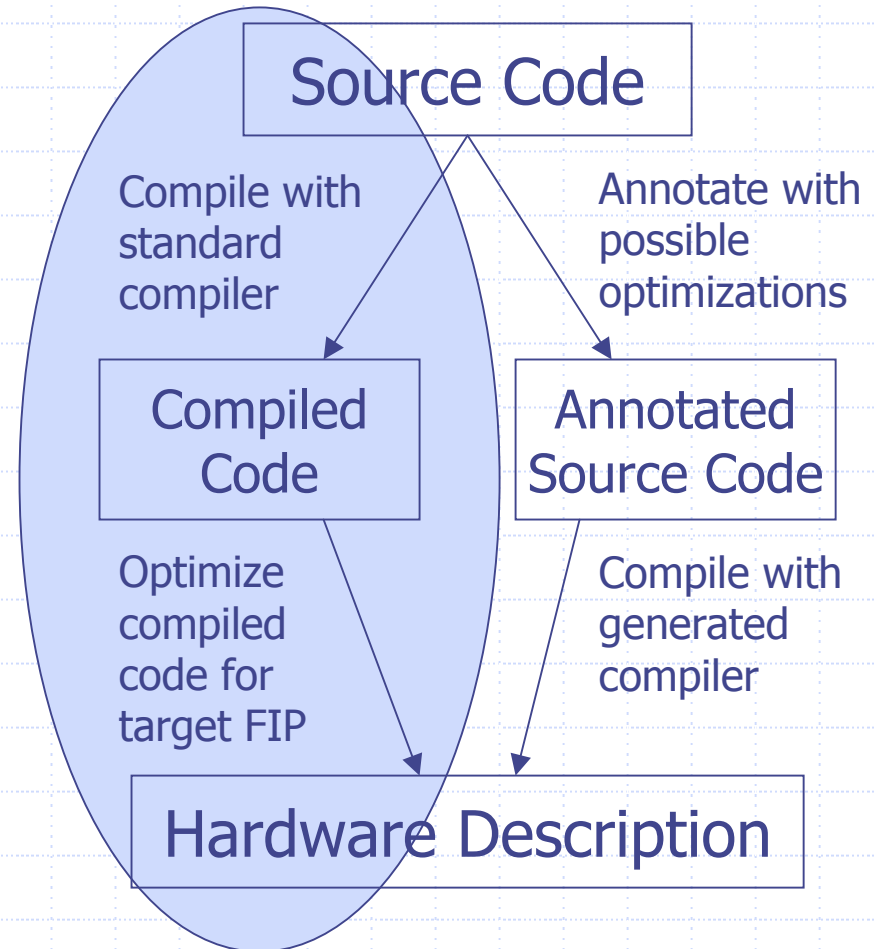
Enhancing Processor Templates



- ◆ Can group modules differently to achieve different processor styles
- ◆ Superscalar
 - More than one execution module
- ◆ Hybrid
 - Different execution modules (stack/register)

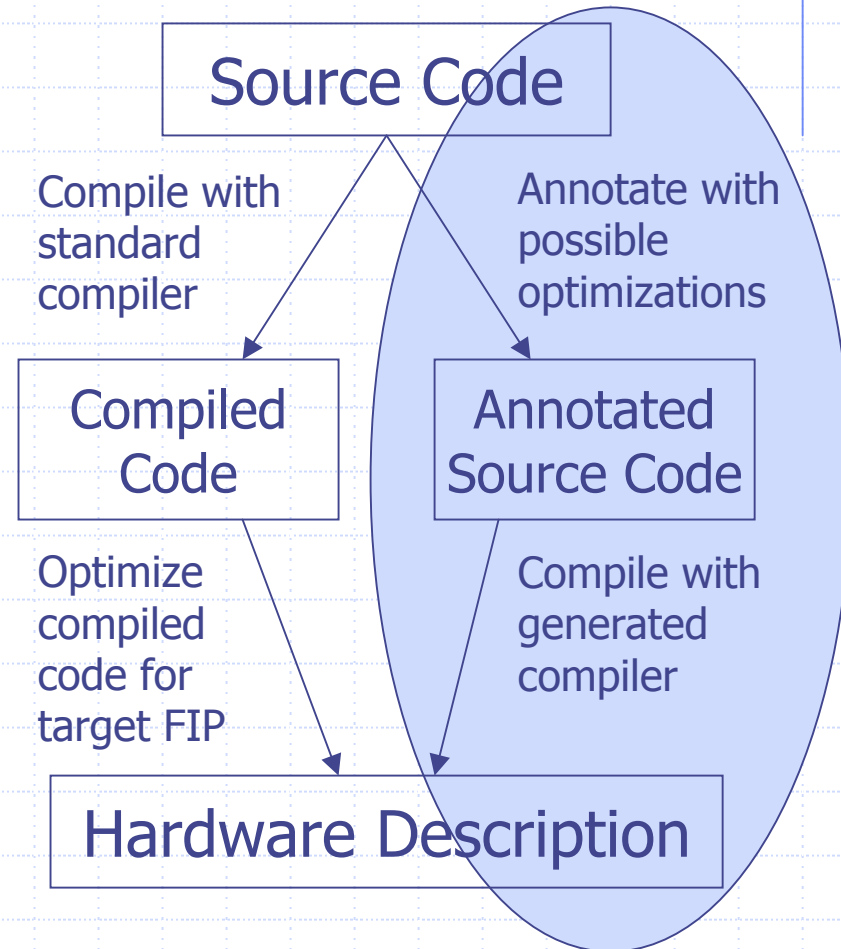
Compilation Strategy

- ◆ Optimizes known target code to known target FIP
- ◆ Reorganizes code to take advantage of ILP
- ◆ Similar to Java's Just-In-Time compilation



Compilation Strategy

- ◆ Annotate source code with profile information
 - Frequency of instruction use
 - Common instruction groups
 - Shared resources
- ◆ Use annotations to drive optimization of code and FIP
- ◆ Compile with a FIP-specific compiler



FIP Optimizations

◆ Custom instructions

- Direct hardware implementations can be exposed by custom instructions
- Reduces number of fetch-decode instructions
- Increases area

◆ Processor-specific


- Can replicate much-used resources
- JVM: allow concurrent reads of top of stack
- JVM: allow reads of top 'n' elements

Implementation: JVM

◆ Implemented FIP-based Java Virtual Machine

- #1: segregated resources that are shared
 - ◆ Good area utilization and the expense of routing congestion
- #2: two-stage pipeline & only shares stack and main memory
- #3: deeper pipelines for certain instructions and 'register' style improvements
 - ◆ Top-of-stack registers are replicated – writes are done back to stack directly, replicated registers updated during the fetch cycle
 - ◆ Most instructions have 4 stages, some 5-6 stages
 - ◆ Routing is pipelined

Results



Results

Devices	CaffineMark	Speedup
Software JVM	502	1
FIP JVM	1019	2x
Pipelined FIP JVM	3665	7x
GMJ30501SB	13332	27x

FIP implemented in Xilinx Virtex @33MHz
GMJ30501SB == 200MHz ASIC

Impact of custom instructions

◆ Link list insertion (100 links)

- Direct hardware: 2.3ms
- FIP-JVM: 6.4ms, 22 instructions
- ASIC: 1ms

◆ Adding 'insert link' instruction

- 12 cycles to complete
- Saves 10 cycles and 10 fetch/decodes

Implementation: AES

Device	Cycles/Block	Slices/BRAM	Mbps/MHz
Software	340	N/A	0.4
FIP (I)	99	1770/2	1.3
FIP (II)	32	1393/10	4
Hardware I	11	460/10	11.5
Hardware II	1	2679/82	129.6

FIP (I) direct implementation of transforms

FIP (II) table-lookups for transforms

Hardware I Spartan II 100-6

Hardware II Virtex-E 812-8

Optimizations

◆ Opcode chaining

- Store intermediate results in temporary registers
- Removes push/pop and alleviates sequential nature of stack-based machines

◆ Galois Multiplication is single instruction

- No procedure call overhead
- 8.5 cycles instead of 30

◆ Introduce predicate registers to allow instructions to execute in parallel

- 6 cycles instead of 30